

# What is Wrong with Topic Modeling? (and How to Fix it Using Search-based SE)

Amritanshu Agrawal\*, Wei Fu\*, Tim Menzies\*

Department of Computer Science, North Carolina State University, Raleigh, NC, USA

---

## Abstract

**Context:** Topic modeling finds human-readable structures in unstructured textual data. A widely used topic modeler is Latent Dirichlet Allocation. When run on different datasets, LDA suffers from “order effects” i.e. different topics are generated if the order of training data is shuffled. Such order effects introduce a systematic error for any study. This error can relate to misleading results; specifically, inaccurate topic descriptions and a reduction in the efficacy of text mining classification results.

**Objective:** To provide a method in which distributions generated by LDA are more stable and can be used for further analysis.

**Method:** We use LDADE, a search-based software engineering tool that tunes LDA’s parameters using DE (Differential Evolution).

**Evaluation:** LDADE is evaluated on data from a programmer information exchange site (StackOverflow), title and abstract text of thousands of Software Engineering (SE) papers, and software defect reports from NASA. Results were collected across different implementations of LDA (Python+Scikit-Learn, Scala+Spark); across different platforms (Linux, Macintosh) and for different kinds of LDAs (VEM, or using Gibbs sampling). Results were scored via topic stability and text mining classification accuracy.

**Results:** In all treatments: (i) standard LDA exhibits very large topic instability; (ii) LDADE’s tunings dramatically reduces cluster instability; (iii) LDADE also lead to improved performances for supervised as well as unsupervised learning.

**Conclusion:** Due to topic instability, using standard LDA with its “off-the-shelf” settings should now be depreciated. Also, in future, we should require SE papers that use LDA to test and (if needed) mitigate LDA topic instability. Finally, LDADE is a candidate technology for effectively and efficiently reducing that instability.

**Keywords:** Topic modeling, Stability, LDA, tuning, differential evolution.

---

## 1. Introduction

The current great challenge in software analytics is understanding unstructured data. As shown in Figure 1, most of the planet’s 1600 Exabytes of data does not appear in structured sources (databases, etc) [48]. Rather it is *unstructured* data, often in free text, and found in word processing files, slide presentations, comments, etc etc.

Such unstructured data does not have a pre-defined data model and is typically text-heavy. Finding insights among unstructured text is difficult unless we can search, characterize, and classify the textual data in a meaningful way. One of the common techniques for finding related topics within unstructured text (an area called topic modeling) is Latent Dirichlet Allocation (LDA) [7].

This paper explores systematic errors in LDA analysis. LDA is a non-deterministic algorithm since its internal weights are updated via a stochastic sampling process (described later in this paper). We show in this paper that this non-determinism means that the topics generated by LDA on SE data are subject to order effects; i.e. different input orderings can lead to different topics. Such instability can:

- Confuse users when they see different topics each time the algorithm is re-run)
- Reduce the efficacy of text mining classifiers that rely on LDA to generate their input training data.

To fix this problem, we propose LDADE: a combination of LDA and a search-based optimizer (differential evolution, or DE) [61]) that automatically tunes LDA’s  $\langle k, \alpha, \beta \rangle$  parameters. This paper tests LDADE by applying text mining to three data sets:

1. Data from a programmer information exchange site (Stackoverflow);
2. Title and abstract text of 15121 SE papers;
3. Software defect reports from NASA.

Using these datasets, we explore these research questions:

- **RQ1: Are the default settings of LDA incorrect?** We will show that using the default settings of LDA for SE data can lead to systematic errors since stability scores start to drop after  $n = 5$  terms per topic.
- **RQ2: Does LDADE improve the stability scores?** LDADE dramatically improves stability scores using the parameters found automatically by DE.
- **RQ3: Does LDADE improve text mining classification accuracy?** Our experiments shows that LDADE also improves classification accuracy.

---

\*Corresponding author: Tel: +1-919-637-0412 (Amritanshu)

Email addresses: aagrawa8@ncsu.edu (Amritanshu Agrawal), wfu@ncsu.edu (Wei Fu), tim.menzies@gmail.com (Tim Menzies)

URL: <http://ai4se.net/> (Amritanshu Agrawal)




Figure 1: Data Growth 2005-2015. From [48].

- **RQ4: Do different data sets need different configurations to make LDA stable?** LDADE finds different “best” parameter settings for different data sets. Hence reusing tunings suggested by other researchers from other data sets is *not* recommended. Instead, it is better to use automatic tuning methods to find the best tuning parameters for the current data set.
- **RQ5: Using different kinds of LDA or with different implementations, are our findings consistent?** Results were collected across different implementations of LDA (Python+Scikit-Learn, Scala+Spark); across different platforms (Linux, Macintosh) and for different kinds of LDAs (the traditional VEM method, or using Gibbs sampling). Across all these implementations, the same effect holds: (a) standard LDA suffers from order effect and topic instability and (b) LDADE can reduce that instability.
- **RQ6: Is tuning easy?** We show that, measured in the terms of the internal search space of the optimizer, tuning LDA is much simpler than standard optimization methods.
- **RQ7: What is the runtime cost of tuning?** The advantages of LDADE come at some cost: tuning with DE makes LDA three to five times slower. While this is definitely more than not using LDA, this may not be an arduous increase given modern cloud computing environments.
- **RQ8: Should topic modeling be used “off-the-shelf” with their default tunings?** Based on these findings, our answer to this question is an emphatic “no”. We can see little reason to use “off-the-shelf” LDA for any kind of SE data mining applications.

The rest of this paper is structured as follows. §2 argues that stabilizing the topics generated by LDA is important for several reasons:

1. LDA is widely used so it is important to study how to use it better;
2. A lack of stability in LDA topics raises doubts about some prior research results;
3. Improving LDA topic stability improves classification accuracy;

Related work is revised in §3 and the methods of this paper are discussed in §4. The above research questions are answered in §5. This is followed by a discussion on the validity of our results and a section describing our conclusions. Note that the main conclusion of this paper is that, henceforth, we should require SE papers that use LDA to test and (if needed) mitigate LDA topic instability.

## 2. Motivation

### 2.1. LDA is Widely Used

We study LDA since this algorithm is a widely-used technique in recent research papers appearing in prominent SE venues. Table 1 and 2 shows SE venues that publish SE results as well as a sample of those papers.

As witnessed by the central columns of Table 2, many prior papers have commented that the results of a topic modeling analysis can be affected by tuning the control parameters of LDA. Yet as reported in §3, a repeated pattern in the literature is that, despite these stated concerns, researchers rarely take the next step to find ways to find better control tunings.

### 2.2. Standard LDA Can Make Misleading Conclusions

Standard practice in papers that use LDA is to present a table showing the top, say, 40 topics. This section shows one example where, due to LDA instability, the contents of such tables can only be described as mostly illusionary. Later in this paper (in §5.1) we show that this example is actually representative of a general problem: changing the input ordering dramatically changes the topics reported by standard LDA.

The example of this section comes from a recent analysis of topics and trends in the StackOverflow programmer discussion forum. Barua et al. [2] analysed topics and trends in StackOverflow. In their paper, they listed the top 40 topics found at StackOverflow.

Venue	Full Name	Count
ICSE	International Conference on Software Engineering	4
CSMR-WCRE / SANER	International Conference on Software Maintenance, Reengineering, and Reverse Engineering / International Conference on Software Analysis, Evolution, and Reengineering	3
ICSM / ICSME	International Conference on Software Maintenance / International Conference on Software Maintenance and Evolution	3
ICPC	International Conference on Program Comprehension	4
ASE	International Conference on Automated Software Engineering	3
ISSRE	International Symposium on Software Reliability Engineering	2
MSR	International Working Conference on Mining Software Repositories	8
OOPSLA	International Conference on Object-Oriented Programming, Systems, Languages, and Applications	1
FSE/ESEC	International Symposium on the Foundations of Software Engineering / European Software Engineering Conference	1
TSE	IEEE Transaction on Software Engineering	1
IST	Information and Software Technology	3
SCP	Science of Computer Programming	2
ESE	Empirical Software Engineering	4

Table 1: SE Venues that publish on Topic Modeling.

ID	Year	Citations	Venues	Mentions instability in LDA?	Uses Default Parameters	Does tuning?	Conclusion	Tasks / Use cases
[59]	2011	112	WCRE	Y	Y	N	Explored Configurations without any explanation.	Bug Localisation
[51]	2010	108	MSR	Y	Y	N	Explored Configurations without any explanation. Reported their results using multiple experiments.	Traceability Link recovery
[2]	2014	96	ESE	Y	Y	N	Explored Configurations without any explanation. Choosing right set of parameters is a difficult task.	StackOverflow Q&A data analysis
[55]	2013	75	ICSE	Y	Y	Y	Uses GA to tune parameters. They determine the near-optimal configuration for LDA in the context of only some important SE tasks.	Finding near-optimal configurations
[17]	2013	61	ICSE	Y	Y	N	Explored Configurations without any explanation.	Software Requirements Analysis
[26]	2011	45	MSR	Y	Y	N	They validated the topic labelling techniques using multiple experiments.	Software Artifacts Analysis
[24]	2014	44	RE	Y	Y	N	Explored Configurations without any explanation.	Requirements Engineering
[64]	2011	44	ICSE	Y	Y	N	Open issue to choose optimal parameters.	A review on LDA
[65]	2014	35	SCP	Y	Y	N	Explored Configurations without any explanation.	Software Artifacts Analysis
[10]	2012	35	MSR	Y	Y	N	Choosing the optimal number of topics is difficult.	Software Defects Prediction
[66]	2014	31	ESE	Y	Y	N	Choosing right set of parameters is a difficult task.	Software Testing
[1]	2009	29	MSR	Y	Y	N	Explored Configurations without any explanation and accepted to the fact their results were better because of the corpus they used.	Software History Comprehension
[36]	2013	27	ESEC/FSE	Y	Y	Y	Explored Configurations using LDA-GA.	Traceability Link recovery
[5]	2014	20	ICPC	Y	Y	N	Use heuristics to find right set of parameters.	Source Code Comprehension
[35]	2013	20	MSR	Y	Y	N	In Future, they planned to use LDA-GA.	StackOverflow Q&A data analysis
[27]	2014	15	WebSci	Y	Y	N	Explored Configurations without any explanation.	Social Software Engineering
[22]	2013	13	SCP	Y	Y	N	Their work focused on optimizing LDAs topic count parameter.	Source Code Comprehension
[25]	2012	13	ICSM	Y	Y	N	Explored Configurations without any explanation.	Software Requirements Analysis
[16]	2015	6	IST	Y	Y	N	Explored Configurations without any explanation. Choosing right set of parameters is a difficult task.	Software re-factoring
[18]	2016	5	CS Review	Y	Y	N	Explored Configurations without any explanation.	Bibliometrics and citations analysis
[33]	2014	5	ISSRE	N	Y	N	Explored Configurations without any explanation.	Bug Localisation
[50]	2015	3	JIS	Y	Y	N	They improvised LDA into ISLDA which gave stability across different runs.	Social Software Engineering
[62]	2015	2	IST	Y	Y	Y	Explored Configurations using LDA-GA.	Software Artifacts Analysis
[9]	2016	0	JSS	N	Y	N	Explored Configurations without any explanation. Choosing right set of parameters is a difficult task.	Software Defects Prediction

Table 2: A sample of the recent literature on using topic modeling in SE. Note that some of these papers are widely-cited.

Figure 2 shows our attempt to reproduce their results. Note that we could not produce a verbatim reproduction since they used a data dump from StackOverflow using data available before 2012 while all we could access was current data. In that figure:

- We ran LDA twice with different randomly generated input orderings (keeping all other parameter settings intact);
- The topics from the first run were scored and sorted according to the overlap of their words from the second run.

Figure 2 shows the results. While a very few topics appear verbatim in the two runs (see the top three lines of Figure 2), most do not. In fact, 25 of the 40 topics in that figure show instability (have an overlap of 55% or less).

Figure 2 shows the results where we only examine the top  $n = 9$  words in each topic of run1 and run2. We selected  $n = 9$  since, later in this paper, we find that an analysis of  $n > 9$  words per topics leads to near zero percent overlap of the topics generated in different runs (see §5.1). That is, the instabilities of Figure 2 get even *worse* if we use *more* of the LDA output.

### 2.3. LDA Stabilization Means Better Inference

Inference with LDA can be assessed:

- Via topic similarities (as done in Figure 2);
- Via the classification performance if the LDA topics are used as features to be fed into a classifier.

As shown later in this paper, we can use LDADE to increase the similarities of the LDA topics generated by LDA (see §5.2).

As to the effects on classification accuracy, one way to score a classifier is via the  $F_\beta$  score that combines *precision*  $p$  and *recall*  $r$  as follows:

$$F_\beta = (1 + \beta^2) \frac{pr}{p\beta^2 + r} \quad (1)$$

In our lab, we work closely with industrial partners who build text mining classifiers to report if legal documents are “relevant” or “irrelevant” to a current case. These industrial partners score their classifiers using  $F_2$ ; i.e.  $5pr/(4p + r)$ . This  $F_2$  score is useful when reporting classification results since it favors classifiers that do not waste the time of industrial clients with false positives.

% overlap with closest topic in run2	Topic name	top 9 words in topic
100	Xaml Binding	grid window bind valu wpf silverlight control xaml properti
100	Ruby on Rails	rubi rail gem lib user end rvm app requir
100	MVC	http com java apach bean org springframework sun servlet
88	Objective C	self cell nsstring iphon nil object anim alloc view
88	Function Return Types	function const char void includ amp return int std
77	Testing	chang tri work use problem like set code test
77	Socket Communications	Main send socket connect run Android Activity thread start task process time
77	Q and A	know need want way use question time like make
77	OO Programming	instanc void type new method class return object public
77	HTML Links	com http www url html site content page link
77	Display	height jpg imag png src size img color width
66	Windows/VS Tips	studio net dll web use mvc control asp visual
66	Website Design	left span style text class div color css width
66	Web Development	script jquery function form input type ajax var javascript
66	Java Programming	java void new privat return null int public string
66	Date/Time Format	select day valu option year format month time date
66	Database	column order group select join tabl row queri null
66	Android User View	height content textview parent wrap android width view layout
66	.NET Framework	net form text checkbox control click label asp button
55	iOS App Development	applic user need iphon develop want app use like
55	MySQL	connect sql databas tabl row queri data mysql server
55	HTML Form	form login user usernam password page XML control action session
55	Git Operations	folder upload open git use path read file directori
55	Email Message	address form send field email messag error valid contact
55	Eclipse	maven jar target build eclips version depend plugin project
44	iOS App Development	applic user need iphon develop want app use like
44	Regular Expressions	function array replac regex match echo php post string
44	Media Player	object video play game player obj use data audio
44	Float number Manipulation	doubl valu count rang number data length float point
44	Compiling	librari compil lib includ usr command file error test
44	.NET Framework	net form text checkbox control click label asp button
33	Ruby Version Manager	end lib user rvm rubi rail app gem requir
33	Scripting Language	def line self modul print import templat django python
33	NodeJS	tag function node express parent use list like variabl
33	C# Programming	net web asp servic mvc use control applic wcf
33	Python Programming	line python file print command script curl output run
33	Android Debugging	com java lang debug method info android error androidruntime
22	Visual Studio	window project file error librari compil visual build dll
22	Information Systems	messag email log address contact send mail phone locat
22	Flex4 Development	flash function new var list click sub event item

Figure 2: LDA topic instability. Shows results from two runs. Column1 reports the number of words from a topic seen in its nearest match in the second run.

§5.3 of this paper compares the  $F_2$  scores seen in text mining classification using standard and stable LDA topics. There, we report significant and large  $F_2$  improvements. That is, LDADE not only improves topic stability, but also the efficacy of the inference that subsequently uses those topics.

Topic: String Manipulation	Topic: Function	Topic: OO Programming	Topic: UI Development	Topic: File Operation
string	function	class	control	file
charact	paramet	method	view	director
encod	pass	object	event	path
format	return	call	button	folder
convert	argument	interfac	click	creat

Figure 3: Example (from [21]) of generating topics from StackOverflow. For each topic, we show just the five most heavily weighted words.

### 3. Related Work

#### 3.1. Topic Modeling

Latent Dirichlet Allocation (LDA) is a generative statistical model that allows sets of observations to be explained by unobserved groups that explain why some parts of the data are similar. It learns the various distributions (the set of topics, their associated word probabilities, the topic of each word, and the particular topic mixture of each document). What makes topic modeling interesting is that these algorithms scale to very large text corpuses. For example, in this paper, we apply LDA to whole of StackOverflow, as well as to two other large text corpuses in SE.

Figure 3 illustrates topic generation from StackOverflow. To find these topics, LDA explores two probability distributions:

- $\alpha = P(k|d)$ ; probability of topic  $k$  in document  $d$ .

- $\beta = P(w|k)$ ; probability of word  $w$  in topic  $k$ ;

Initially,  $\alpha, \beta$  may be set randomly as follows: each word in a document was generated by first randomly picking a topic (from the documents distribution of topics) and then randomly picking a word (from the topics distribution of words). Successive iterations of the algorithm count the implications of prior sampling which, in turn, incrementally updates  $\alpha, \beta$ .

Binkley et al. [5] performed an extensive study and found that apart from  $\alpha, \beta$ , the other parameters that define LDA are:

- $k$  = number of topics;
- $b$  = number of burn-in iterations;
- $si$  = the sampling interval

Binkley et al.’s study of the LDA settings was a mostly manual process guided by their considerable background knowledge and expertise and program comprehension. In the field of program comprehension, the Binkley article is the state of the art in applications of LDA to software engineering.

To that work, this paper adds a few extra conclusions. Firstly, we explore LDA in fields other than program comprehension. Secondly, we ask the question “what if the analysts lacks extensive background knowledge of the domain?”. In that circumstance, some automatic method is needed to support an informed selection of the LDA parameters.

### 3.2. About Order Effects

This paper uses tuning to fix “order effects” in topic modeling. Langley [19] defines such effects as follows:

*A learner  $L$  exhibits an order effect on a training set  $T$  if there exist two or more orders of  $T$  for which  $L$  produces different knowledge structures.*

Many learners exhibit order effects: e.g. certain incremental clustering algorithms generate different clusters, depending on the order with which they explore the data [19]. Hence, some algorithms survey the space of possible models across numerous random divisions of the data (e.g. Random Forests [8]).

From the description offered above in §3.1, we can see (a) how topic modeling might be susceptible to order effects and (b) how such order effects might be tamed:

- In the above description,  $k$ ,  $\alpha$ ,  $\beta$  are initialized at random then updated via an incremental re-sampling process. Such incremental updates are prone to order effects.
- One technique to reduce the effect of different data orderings is to initialize,  $k$ ,  $\alpha$ ,  $\beta$  to some useful value. As shown in §5.4, the key to applying this technique is that different data sets will requiring different initializations; i.e. the tuning process will have to be repeated for each new data set.

### 3.3. Tuning: Important and (Mostly) Ignored

The impact of tuning are well understood in the theoretical machine learning literature [4]. When we tune a data miner, what we are really doing is changing how a learner applies its heuristics. This means tuned data miners use different heuristics, which means they ignore different possible models, which means they return different models; i.e. *how* we learn changes *what* we learn.

Yet issues relating to tuning are poorly addressed in the software analytics literature. Fu et al [15] surveyed hundreds of recent SE papers in the area of software defect prediction from static code attributes. They found that most SE authors do not take steps to explore tunings (rare exception: [63]). For example, Elish et al [12] compared support vector machines to other data miners for the purposes of defect prediction. That paper tested different “off-the-shelf” data miners on the same data set, without adjusting the parameters of each individual learner. Similar comparisons of data miners in SE, with no or

minimal pre-tuning study, can be found in the work on Lessmann et al. [34] and, most recently, in Yang et al. [73].

In summary, all our literature reviews of the general (non-LDA) software analytics literature show that the importance of tuning is often mentioned, but never directly address.

### 3.4. LDA and Instability and Tuning

Within the LDA literature, some researchers have explored LDA instability. In April 2016, we searched scholar.google.com for the conjunction of “lda” and “topics” or “stable” or “unstable” or “coherence”. Since 2012, there are 189 such papers, 57 of which are related to software engineering results. Table 2 gives a broad discussion on these papers. In short, of those papers:

- 29/57 mention instability in LDA.
- Of those 29, despite mentioning stability problems, 10 papers still used LDA’s “off-the-shelf” parameters;
- The other 29-10=19 papers used some combination of manual adjustment or some under-explained limited exploration of tunings based on “engineering judgment” (i.e. some settings guided by the insights of the researchers).
- Only 4 of the authors acknowledge that tuning might have a large impact on the results.

Apart from tuning, there are several other workarounds explored in the literature in order to handle LDA instability. Overall, there was little systematic exploration of tuning and LDA in the SE literature. Instead, researchers relied on other methods that are less suited to automatic reproduction of prior results.

In the literature, researchers [24, 38, 39] manually accessed the topics and then used for further experiments. Some made use of Amazon Mechanical Turk to create gold-standard coherence judgements [30]. All these solutions are related to results stability rather than model stability. Note that this workaround takes extensive manual effort and time.

Another approach to tame LDA instability is to incorporate user knowledge into the corpus. For example, SC-LDA [72] can handle different kinds of knowledge such as word correlation, document correlation, document label and so on. Using such user knowledge, while certainly valuable, is somewhat subjective. Hence, for reasons of reproducibility, we prefer fully automated methods.

Some researchers used genetic algorithms to learn better settings for LDA [36, 55, 62]. Genetic algorithms are themselves a stochastic search process. That is, the changes to input orderings explored in this paper would introduce further conclusion instability from the genetic algorithms. In principle, that instability could be removed via extra runs of genetic algorithms over multiple sub-samples of that, where the GA goals are augmented to include “similar topics should be found in different runs”. That said:

- None of the prior work using GAs to improve LDA have applied those sub-sampling stability test;
- If done naively, adding further goals and data sub-sampling to a GA runs the risk of dramatically increasing the runtimes. One reason to prefer LDADE is that it terminates very quickly.

Data set	Size	
	Before Preprocessing	After Preprocessing
PitsA	1.2 MB	292 KB
PitsB	704 KB	188 KB
PitsC	143 KB	37 KB
PitsD	107 KB	26 KB
PitsE	650 KB	216 KB
PitsF	549 KB	217 KB
Citemap	8.6 MB	3.7 MB
StackOverflow	7 GB	589 MB

Table 3: Statistics on our datasets. PitsA, PitsB, etc refer to the issues from six different NASA projects.

Finally, other researchers explore some limited manual parameter tuning for LDA (e.g. experiment with one parameter: cluster size) [17, 67] achieved higher stability by just increasing the number of cluster size. Note that the automatic tuning methods explored by this paper can explore multiple parameters. Further, our analysis is repeatable.

## 4. Methods

This section describes our evaluation methods for measuring instability as well as the optimization methods used to reduce that instability.

### 4.1. Data Sets

To answer our research questions, and to enable reproducibility of our results, we use three open source datasets summarized in Table 3 and described below.

**PITS** is a text mining data set generated from NASA software project and issue tracking system (PITS) reports [41, 44]. This text discusses bugs and changes found in big reports and review patches. Such issues are used to manage quality assurance, to support communication between developers. Topic modeling in PITS can be used to identify the top topics which can identify each severity separately. The dataset can be downloaded from the PROMISE repository [43]. Note that, this data comes from six different NASA projects, which we label as PitsA, PitsB, etc.

**StackOverflow** is the flagship site of the Stack Exchange Network which features questions and answers on a wide range of topics in computer programming. Topic modeling on StackOverflow is useful for finding patterns in programmer knowledge. This data can be downloaded online<sup>1</sup>.

**Citemap** contains titles and abstracts of 15121 papers from a database of 11 senior software engineering conferences from 1992-2016. Most of this data was obtained in the form of an SQL dump from the work of Vasilescu et al. [68]. This dataset is available on-line<sup>2</sup>.

For this study, all datasets were preprocessed using the usual text mining filters [14]:

- Stop words removal using NLTK toolkit<sup>3</sup> [6] : ignore very

common short words such as “and” or “the”.

- Porter’s stemming filter [57]: delete uninformative word endings; e.g. after stemming, all the following words would be rewritten to “connect”: “connection”, “connections”, “connective”, “connected”, “connecting”.
- Tf-idf feature selection: focus on the 5% of words that occur frequently, but only in small numbers of documents. If a word occurs  $w$  times and is found in  $d$  documents and there are  $W, D$  total number of words and documents respectively, then tf-idf is scored as follows:

$$tfidf(w, d) = \frac{w}{W} * \log \frac{D}{d}$$

Table 3 shows the sizes of our data before and after preprocessing. These datasets are of different sizes and so are processed using different tools:

- PITS and Citemap is small enough to process on a single (four core) desktop machine using Scikit-Learn [56] and Python.
- StackOverflow is so large (7GB) that its processing requires extra hardware support. This study used Spark and Mlib on a cluster of 45 nodes to reduce the runtime.

### 4.2. Similarity Scoring

To evaluate topics coherence in LDA, there is a direct approach, by asking people about topics, and an indirect approach by evaluating *pointwise mutual information (PMI)* [30, 53] between the topic words. We could not use any of these criteria, as it requires experts to have domain knowledge. *Perplexity* is the inverse of the geometric mean per-word likelihood. The smaller the perplexity, the better (less uniform) is the LDA model. The usual trend is that as the value of perplexity drops, the number of topics should grow [27]. Researchers caution that the value of perplexity doesn’t remain constant with different topic size and with dictionary sizes [74]. A lot depend on the code implementation of perplexity and the type of datasets used. Since, we are using different implementations of LDA across different platforms on various datasets, we are not using perplexity as evaluation measure.

Another approach researchers have used is *Jaccard Similarity* [17, 53]. We define our measure similar to it. For this work, we assess topic model stability via the *median number overlaps of size  $n$  words (size\_of\_topic)*, which we denote  $\mathfrak{R}_n$ .

For this measurement, we first determine the maximum size of topics we will study. For that purpose, we will study the case of  $n \leq 9$  (we use 9 as our maximum size since the cognitive science literature tells us that  $7 \pm 2$  is a useful upper size for artifacts to be browsed by humans [45]).

Next, for  $1 \leq n \leq 9$ , we will calculate the median size of the overlap, computed as follows:

- Let one *run* of our rig shuffle the order of the training data, then builds topic models using the data;
- $m$  runs of our rig executes  $m$  copies of one run, each time using a different random number seed,
- We say topics are stable, when there are  $x$  occurrences of  $n$  terms appearing in all the topics seen in the  $m$  runs.

<sup>1</sup><http://tiny.cc/SOProcess>

<sup>2</sup><https://github.com/ai-se/citemap/blob/master/citemap.csv>

<sup>3</sup><http://www.nltk.org/book/ch02.html>

RUN 1:	$\mathfrak{R}$ SCORES
Topic 0: glori telemetri command spacecraft trace smrd tim spec pip parent	2/4 = 0.50
Topic 1: spec smrd parent child glori artifact referenc verif matrix data	3/4 = 0.75
Topic 2: test case accuraci roll document glori plan pitch yaw valu	1/4 = 0.25
Topic 3: command specifi ground softwar telemetri initi pip data configur band	4/4 = 1.00

RUN 2:	$\mathfrak{R}$ SCORES
Topic 0: command specifi ground softwar initi telemetri data pip configur band	4/4 = 1.00
Topic 1: document test glori initi plan roll case pitch yaw point	1/4 = 0.25
Topic 2: spec smrd parent child glori artifact verif point matrix spacecraft	3/4 = 0.75
Topic 3: pip capabl glori ground command smrd spec tim list spacecraft	2/4 = 0.50

RUN 3:	$\mathfrak{R}$ SCORES
Topic 0: spec smrd parent child glori referenc artifact verif matrix spacecraft	3/4 = 0.75
Topic 1: command telemetri pip ground spacecraft modesoftware document spec glori	2/4 = 0.50
Topic 2: command specifi ground softwar initi telemetri data configur pip band	4/4 = 1.00
Topic 3: test initi accuraci glori roll plan matrix yaw pitch valu	1/4 = 0.25

RUN 4:	$\mathfrak{R}$ SCORES
Topic 0: command specifi tim pip ground telemetri glori includ softwar document	2/4 = 0.50
Topic 1: test initi glori plan roll accuraci case pip point yaw	1/4 = 0.25
Topic 2: command specifi ground softwar telemetri initi data pip configur band	4/4 = 1.00
Topic 3: spec smrd child glori artifact referenc parent verif matrix data	3/4 = 0.75

Figure 4: Example of topics overlap off size  $n = 5$  across multiple runs.

For example, consider the topics shown in Figure 4. These are generated via four *runs* of our system. In this hypothetical example, we will assume that the runs of Figure 4 were generated by an LDA suffering from topic instability. For  $n = 5$ , we note that Topic 0 of run1 scores  $\frac{2}{4} = 0.5$  since it shares 5 words with topics in only two out of four runs. Repeating that calculation for the other run1 topics shows that:

- Topic 1 of run1 scores  $\frac{3}{4} = 0.75$ ;
- Topic 2 of run1 scores  $\frac{1}{4} = 0.25$ ;
- Topic 3 of run1 scores  $\frac{4}{4} = 1$ .

From this information, we can calculate  $\mathfrak{R}_5$  (the *median number overlaps of size  $n = 5$  words*) as:

$$\text{median}(0.5, 0.75, 0.25, 1) = 0.625$$

Figure 5 shows the  $\mathfrak{R}_n$  scores of Figure 4 for  $1 \leq n \leq 9$ . From this figure, we can see LDA topic instability since any report of the contents of a topic that uses more than three words per topic would be unreliable.

For the following analysis, we distinguish between the **Raw score** and the **Delta score**:

- The two **Raw scores** are the  $\mathfrak{R}_n$  median similarity scores seen *before* and *after* tuning LDA;
- The **Delta score** is the difference between the two **Raw scores** (after tuning - before tuning).

The pseudocode for these calculations is shown in Algorithm 1 with the default set of parameters. In the following description, superscript numbers denote lines in the pseudocode. The data ordering is shuffled every time LDA is run<sup>5</sup>. Data is in the form of term frequency scores of each word per document. Shuffling

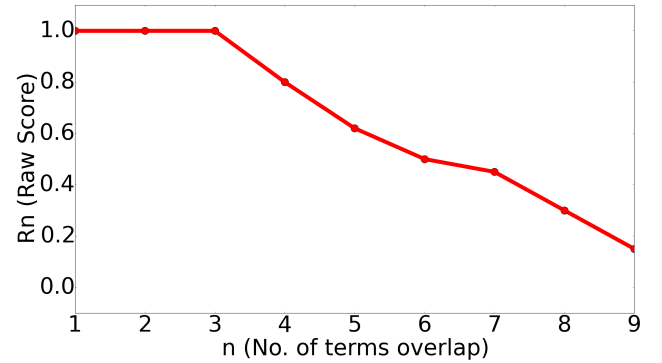


Figure 5:  $\mathfrak{R}_n$  scores of Figure 4 for  $1 \leq n \leq 9$

is done in order to induce maximum variance among the ordering of data with different runs of LDA. Topics<sup>6</sup> are a list of lists which contains topics from all the different runs. A stability score is evaluated on every 10 runs (Fixed), and this process is continued 10 (Fixed) times. At the end, the median score is selected as the untuned raw score ( $\mathfrak{R}_n$ )<sup>3-11</sup>. Hence, the runtimes comes from  $10 * 10$  evaluations of untuned experiment.

#### 4.3. Tuning Topic Modeling with LDADE

LDADE is a combination of topic modeling (with LDA) and an optimizer (differential evolution, or DE) that adjusts the parameters of LDA in order to optimize (i.e. maximize) similarity scores.

---

**Algorithm 1** Pseudocode for untuned LDA with Default Parameters

---

**Input:**  $Data, n, k, \alpha, \beta$  (Defaults)  
**Output:**  $Raw\_Score$

```

1: function LDAScore( $n, Data$ )
2:    $Score \leftarrow \emptyset$ 
3:   for  $j = 0$  to  $10$  do
4:     for  $i = 0$  to  $10$  do
5:        $data \leftarrow \text{shuffle}(Data)$ 
6:        $Topics.add(lda(k, \alpha, \beta))$ 
7:     end for
8:      $Score.add(Overlap(Topics, n, l[0]))$ 
9:   end for
10:   $Raw\_Score \leftarrow \text{median}(Score)$ 
11:  return  $Raw\_Score$ 
12: end function

```

---

We choose to use DE after a literature search on search-based SE methods. The literature mentions many optimizers: simulated annealing [13, 42]; various genetic algorithms [21] augmented by techniques such as DE (differential evolution [61]), tabu search and scatter search [3, 20, 47, 49]; particle swarm optimization [54]; numerous decomposition approaches that use heuristics to decompose the total space into small problems, then apply a response surface methods [28, 75]. Of these, we use DE for two reasons. Firstly, it has proved useful in prior SE tuning studies [71]. Secondly, our reading of the current literature is that there are many advocates for differential evolution.

LDADE adjusts the parameters of Table 4. Most of these parameters were explained above. Apart from them, there are 2 different kinds of LDA implementations as well and they are:

- VEM is the deterministic *variational EM* method that computes  $\alpha, \beta$  via expectation maximization [46].
- Gibbs sampling [23, 71] is a Markov Chain Monte Carlo algorithm, which is an approximate stochastic process for computing and updating  $\alpha, \beta$ . Topic modeling researchers in SE have argued that Gibbs leads to stabler models [31, 32] (a claim which we test, below).

We manually run with these other inference techniques according to different implementations across different platforms. We need to make sure that these instabilities do not hold for just 1 inference technique, or 1 implementation or on 1 platform.

Parameters	Defaults	Tuning Range	Description
$k$	10	[10,100]	Number of topics or cluster size
$\alpha$	None	[0,1]	Prior of document topic distribution. This is called alpha
$\beta$	None	[0,1]	Prior of topic word distribution. This is called beta

Table 4: List of parameters tuned by this paper

Algorithm 2 shows LDADE’s version of DE. DE evolves a *NewGeneration* of candidates from a current *Population*. Each candidate solution in the *Population* is a pair of (Tunings, Scores). Tunings are selected from Table 4 and Scores come similarly from Algorithm 1<sup>3–11</sup>. Note that there is not any outer loop<sup>3</sup>: in Algorithm 2, LDA is only run as 1 rig<sup>11&12</sup>. Here, the runtimes comes from  $iter * np * 10$  evaluations of tuned experiment.

The main loop of DE<sup>9</sup> runs over the *Population*, replacing old items with new Candidates (if new candidate is better). DE generates *new Candidates* via extrapolating<sup>23</sup> between current solutions in the frontier. Three solutions a, b, c are selected at random. For each tuning parameter  $i$ , at some probability  $cr$ , we replace the old tuning  $x_i$  with  $y_i$ . For booleans, we use  $y_i = x_i$  (see line 31). For numerics,  $y_i = a_i + f \times (b_i - c_i)$  where  $f$  is a parameter controlling crossover. The trim function<sup>33</sup> limits the new value to the legal range min..max of that parameter.

The loop invariant of DE is that, after the zero-th iteration<sup>7</sup>, the *Population* contains examples that are better than at least one other candidate. As the looping progresses, the *Population* is full of increasingly more valuable solutions which, in turn, also improves the candidates, which are Extrapolated from the *Population*. Hence, Vesterstrom et al [69] found DE to be competitive with particle swarm optimization and other GAs.

Note that DEs have been applied before for parameter tuning (e.g. see [11, 15, 36, 52, 55, 62]) but this is the first time they have been applied to tune LDA to increase stability.

---

**Algorithm 2** Pseudocode for DE with a constant number of iterations

---

**Input:**  $np=10, f=0.7, cr=0.3, iter=3, Goal \in \text{Finding maximum score}$   
**Output:**  $Raw\_Score, final\_generation$

```

1: function DE( $np, f, cr, iter, Goal$ )
2:    $Cur\_Gen \leftarrow \emptyset$ 
3:    $Population \leftarrow \text{InitializePopulation}(np)$ 
4:   for  $i = 0$  to  $np - 1$  do
5:      $Cur\_Gen.add(Population[i], ldascore(Population[i], n, Data))$ 
6:   end for
7:   for  $i = 0$  to  $iter$  do
8:      $NewGeneration \leftarrow \emptyset$ 
9:     for  $j = 0$  to  $np - 1$  do
10:       $S_i \leftarrow \text{Extrapolate}(Population[j], Population, cr, f, np)$ 
11:      if  $ldascore(S_i) \geq Cur\_Gen[j][1]$  then
12:         $NewGeneration.add(S_i, ldascore(S_i, n, Data))$ 
13:      else
14:         $NewGeneration.add(Cur\_Gen[j])$ 
15:      end if
16:    end for
17:     $Cur\_Gen \leftarrow NewGeneration$ 
18:  end for
19:   $Raw\_Score \leftarrow \text{GetBestSolution}(Cur\_Gen)$ 
20:   $final\_generation \leftarrow Cur\_Gen$ 
21:  return  $Raw\_Score, final\_generation$ 
22: end function
23: function EXTRAPOLATE( $old, pop, cr, f, np$ )
24:   $a, b, c \leftarrow \text{threeOthers}(pop, old)$ 
25:   $newf \leftarrow \emptyset$ 
26:  for  $i = 0$  to  $np - 1$  do
27:    if  $cr \leq \text{random}()$  then
28:       $newf.add(old[i])$ 
29:    else
30:      if  $\text{typeof}(old[i]) = \text{bool}$  then then
31:         $newf.add(not old[i])$ 
32:      else
33:         $newf.add(trim(i, a[i] + f * (b[i] - c[i])))$ 
34:      end if
35:    end if
36:  end for
37:  return  $newf$ 
38: end function

```

---



## 5. Experimentation

In this section, any result from the smaller data sets (Pits and Citemap) come from Python implementation based on Scikit-Learn running on a single 4-core machine. Also, any results from the larger data (StackOverflow) comes from a Scala implementation based on Mllib [40] running on a 45 node Spark system (8 cores per node).

Note that, for the RQ3, there are some intricate details with classification results. After tuning (Goal is still to maximize the  $\mathfrak{R}_n$  score) and finding the optimal ‘K’, we trained a Linear Kernel SVM classifier using document topic distributions as features just like used by Blei et al [7].

### 5.1. RQ1: Are the default settings of LDA incorrect?

This first research question checks the core premise of this article– that changes in the order of training data dramatically affects the topics learned via LDA. Note that if this is *not true*, then there would be no value-added to this paper.

Figure 6 plots  $n$  vs  $\mathfrak{R}_n$  for untuned LDA. Note that the stability collapses the most after  $n = 5$  words. This means that any report of LDA topics that uses more than five words per topic will be changed, just by changing the order of the inputs. This is a significant result since the standard advice in the LDA papers [37, 55] is to report the top 10 words per topic. As shown in Figure 7a, it would be rare that any such 10 word topic would be found across multiple runs.

#### Result 1

*Using the default settings of LDA for SE data can lead to systematic errors due to topic modeling instability.*

### 5.2. RQ2: Does LDADE improve the stability scores?

Figure 7a and Figure 7b shows the stability improvement generated by tuning. Tuning never has any negative effect (reduces stability) and often has a large positive effect– particular after 5 terms overlap. The largest improvement we was in PitsD dataset which for up to 8 terms overlap was 100% (i.e. was always found in all runs). Overall, after reporting topics of up to 7

words, in the majority case (66%), those topics can be found in models generated using different input orderings. Accordingly, our answer to **RQ2** is:

#### Result 2

*For stable clusters, tuning is strongly recommended for future LDA studies.  $\alpha$  and  $\beta$  matters the most for getting good clusters.*




Figure 7a: After tuning: uses parameters learned by DE.

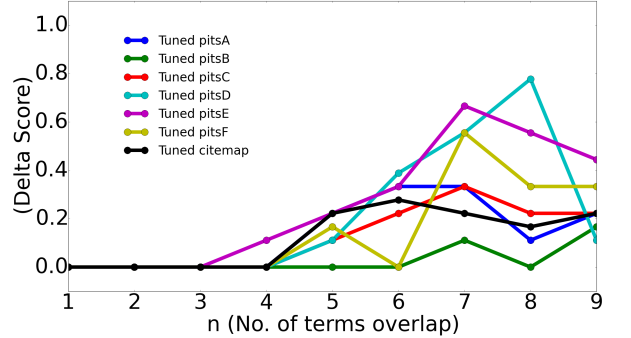


Figure 7b: Delta = After - Before.

Figure 7: **RQ1, RQ2** stability results over ten repeated runs. In these figures, larger numbers are better.




Figure 6: Before tuning: uses LDA’s default parameters

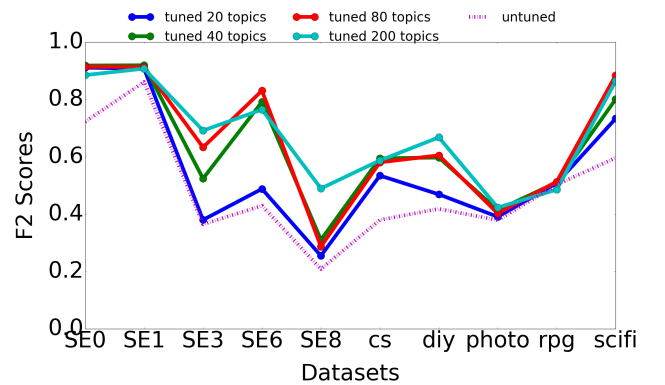


Figure 8: Tuning and Untuned results for Classification SE Task




Figure 9: Datasets vs Parameter ( $k$ ) variation




Figure 10 Datasets vs Parameter ( $\alpha$ ) variation




Figure 11 Datasets vs Parameter ( $\beta$ ) variation

### 5.3. RQ3: Does LDADE improve text mining classification accuracy?

We studied some other StackExchange websites data dump for classification results which were generated by Krishna et al [29]. These datasets are categorized into binary labels saying which documents are relevant and non relevant. Here goal of our DE was still to maximize the  $\mathcal{R}_n$  score. It shouldn't be confused with maximizing some other accuracy goals. After finding the optimal 'K', we trained a Linear Kernel SVM classifier using document topic distributions just like used by Blei et al [7].

In Figure 8, the x-axis represents different datasets as generated. Y-axis represents the F2 score (from Equation 1) which weights recall higher than precision [58]. One legend called as "untuned" uses default parameters with  $k = 10$ . Other legends in the graph shows tuning of  $k$  for 20, 40, 80, 200 by keeping  $\alpha$  and  $\beta$  fixed. There is about 20% minimum improvement than the untuned results.

Hence, we say:

#### Result 3

*For any SE classification task, tuning is again strongly recommended. And  $k$  matters the most for a good classification accuracy.*

### 5.4. RQ4: Do different data sets need different configurations to make LDA stable?

Figures 9, 10, and 11 show the results of tuning for word overlap of  $n = 5$ . On display in each set of vertical bars are the median values generated across 10 tunings. Also, shown are the inter-quartile range (IQR) of those tunings (the IQR is the 75th-25th percentile values and is a non-parametric measure of variation around the median value). Note that in Figure 9, IQR=0 for PitsB dataset where tuning always converged on the same final value.

These figures show how tuning selects the different ranges of parameters. Some of the above numbers are far from the standard values; e.g. Garousi et al [18] recommend using  $k = 67$  topics yet in our data sets, best results were seen using  $k \leq 24$ . Clearly:

#### Result 4

*Do not reuse tunings suggested by other researchers from other data sets. Instead, always re-tune for all new data.*

### 5.5. RQ5: Using different kinds of LDA or with different implementations, are our findings consistent?

Results were collected across different implementations of LDA (Python+Scikit-Learn, Scala+Spark); across different platforms (Linux, Macintosh) and for different kinds of LDAs (the traditional VEM method, or using Gibbs sampling).

To check this, it was insightful to compare our results with:

- The Pits and Citemap results, executed in Scikit-Learn and Python running on a desktop machine.
- The StackOverflow data set executed in Scala using Mllib running on a Spark cluster.

Figure 12 shows tuning results for StackOverflow, Citemap, and PitsA using Scala/Spark cluster (for results on other data sets, see <https://goo.gl/UVaql1>).

Another useful comparison is to change the internal of the LDA:

- Sometimes using VEM sampling;
- Sometimes using Gibbs sampling;

Figure 13 compares the VEM vs Gibbs sampling (for results on other datasets, see <https://goo.gl/faYAcg>).

When compared with the Python/desktop results of Figure 7 we see the same patterns:

- Tuning never makes stability worse.




Figure 12: Spark Results

- Sometimes, it dramatically improves it (in particular, see the Citemap results of Figure 12).

That said, there are some deltas between VEM and Gibbs where it seems tuning is more important for VEM than Gibbs (evidence: the improvements seen after tuning are largest for the VEM results of Figure 13 and at <https://goo.gl/faYAcg>).

#### Result 5

*Instability is not due to any quirk in the implementation of LDA. Instability is consistent and LDADE can stabilize.*

#### 5.6. RQ6: Is tuning easy?

The DE literature recommends using a population size  $np$  that is ten times larger than the number of parameters being optimized [61]. For example, when tuning  $k, \alpha, \beta$ , the DE literature is recommending  $np = 30$ . Figure 14 explores  $np = 30$  vs the  $np = 10$  we use in Algorithm 2 (as well as some other variants of DE's F and CR parameters). The figure shows results just for Citemap and, for space reasons, results relating to other data sets are shown at <https://goo.gl/HQNASF>. After reviewing the results from all the datasets, we can say that there isn't much of an improvement by using different F, CR, and Population size. So our all other experiments used  $F = 0.7$ ,  $CR = 0.3$  and  $np = 10$ . Also:




Figure 13: GIBBS vs VEM

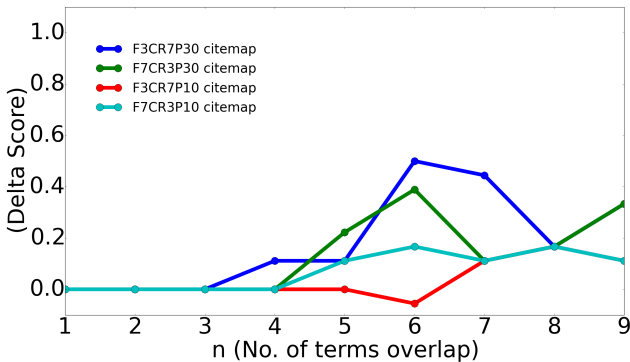


Figure 14: Terms vs Delta Improvement using Different settings of DE

#### Result 6

*Finding stable parameters for topic models is easier than standard optimization tasks.*

#### 5.7. RQ7: What is the runtime cost of tuning?

Search-based SE methods can be very slow. Wang et al [70] once needed 15 years of CPU time to find and verify the tunings required for software clone detectors. Sayyad et al [60] routinely used  $10^6$  evaluations (or more) of their models in order to extract products from highly constrained product lines. Hence, before recommending any search-based method, it is wise to consider the runtime cost of that recommendation.

To understand our timing results, recall that untuned, tuned LDA use Algorithm 1, Algorithm 2 respectively. Based on the pseudocode shown above, our pre-experimental expectation is that tuning will be three times slower than not tuning.

Figures 15 and 16 check if that theoretical holds true in practice. Shown in blue and red are the runtimes required to run LDA untuned, tuned (respectively). The longer runtimes (in red) include the times required for DE to find the tunings. Overall, tuning slows down LDA by a factor of up to five (which is very close to our theoretical prediction). Hence, we say:

#### Result 7

*Theoretically and empirically, tuning LDA costs three to five times more runtime as much as using untuned LDA.*

While this is definitely more than not using DE, but this may not be an arduous increase given modern cloud computing environments.

#### 5.8. RQ8: Should topic modeling be used “off-the-shelf” with their default tunings?

Figure 7 shows that there is much benefit in tuning. Figures 9, 10, and 11 show that the range of “best” tunings is very dataset-specific. Hence, for a new dataset, the off-the-shelf tunings may often fall far from the useful range. Figures 15 and 16 show that tuning is definitely slower than otherwise, but the overall cost is not prohibitive. Hence:

#### Result 8

*Whatever the goal is, whether using the learned topics, or cluster distribution for classification we cannot recommend using “off-the-shelf” LDA.*

## 6. Threats to Validity

As with any empirical study, biases can affect the final results. Therefore, any conclusions made from this work must be considered with the following issues in mind:

**Sampling bias** threatens any experiment; i.e., what matters there may not be true here. For example, the data sets used here




Figure 15: VEM: Datasets vs Runtimes




Figure 16: Gibbs: Datasets vs Runtimes

comes after various pre-processing steps and could change if pre-processed differently. And that is why, the footnotes of this paper contain all the urls where other researchers can download our datasets and explore the conclusions of this paper. Even though we used so many data sets, there could be other datasets for which our results could be wrong or have lesser improvement.

**Learner bias:** For running LDA, we selected other parameters as default which are of not much importance. But there could be some datasets where by tuning them there could be much larger improvement. And for RQ2, we only experimented with linear kernel SVM. There could be other classifiers which can change our conclusions. Data Mining is a large and active field and any single study can only use a small subset of the known data miners.

**Evaluation bias:** This paper uses topic similarity ( $\mathcal{R}_n$ ) and F2 measures of evaluation but there are other measures which are used in software engineering which includes perplexity, performance, accuracy, etc. Assessing the performance of stable LDA is a clear direction for future work.

**Order bias:** With each dataset, how data samples are picked and put into LDA is completely random. Since this paper also consider input order effects, though there could be times when the input order could be with lesser variance. To mitigate this order bias, we ran the experiment 10 times by randomly changing the order of the data samples each time.

Another threat to validity of this work is that it is a quirk of the control parameters used within our DE optimizer. We have some evidence that this is not the case. Figure 14 and <https://goo.gl/HQNASF> explored a range of DE tunings and found little difference across that range. Also, Table V explores another choice within DE – how many evaluations to execute before terminating DEs. All the results in this paper use an evaluation budget of 30 evaluations. Table V compares results across different numbers of evaluations. While clearly, the more evaluations the better, there is little improvement after the 30 evaluations used in this paper.

The conclusions of this paper are based on a finite number of data sets and it is possible that other data might invalidate our

Datasets\Evaluations	10	20	30	50
PitsA	0.9	0.9	1.0	1.0
PitsB	0.9	0.9	0.9	1.0
PitsC	0.9	1.0	1.0	1.0
PitsD	0.9	1.0	1.0	1.0
PitsE	0.9	0.9	1.0	1.0
PitsF	0.9	0.9	0.9	0.9
Citemap	0.67	0.67	0.77	0.77
StackOverflow	0.6	0.7	0.8	0.8

Table 5: Evaluations vs Stability Scores

conclusions. As with all analytics papers, any researcher can do is to make their conclusions and materials public, then encourage other researchers to repeat/ refute/ improve their conclusions. For example, the footnotes of this paper contain all the urls where other researchers can download our materials and explore the conclusions of this paper.

## 7. Conclusion and Future Work

Based on the above, we offer a specific and general conclusion. Most specifically, we recommend

- Any study that shows the topics learned from LDA, and uses them to make a particular conclusion, needs to first tune LDA. We say this since the topics learned from untuned LDA are unstable, i.e. different input orderings will lead to different conclusions. However, after tuning, stability can be greatly increased.
- Unlike the advise of Lukins et al [37], LDA topics should not be reported as the top ten words. Due to order effects, such a report can be highly incorrect. Our results show that up to eight words can be reliably reported, but only after tuning for stability using tools like LDADE.
- Any other studies which are making use of these topic distribution for classification results need to be tuned first before using them in their further tasks.
- Do not download someone else’s pre-tuned LDA since, as shown here, the best LDA tunings vary from data set to data set.

- These results demand us to visit the previous case studies which have used LDA using “off-the-shelf parameters to rerun by tuning these parameters using automated tools like LDADE.

Our experience is that this recommendation is not an arduous demand since tuning adds less than a factor of five to the total run times of an LDA study.

More generally, we comment that the field of software analytics needs to make far more use of search-based software engineering in order to tune their learners. In other work, Fu et al have shown that tuning significantly helps defect prediction [15] (a result that has also been reported by other researchers [55]). In this work, we have shown that tuning significantly helps topic modeling by mitigating a systematic error in LDA (order effects that lead to unstable topics). The implications of this study for other software analytics tasks is now an open and pressing issue. In how many domains can search-based SE dramatically improve software analytics?

## Acknowledgements

The work is partially funded by NSF award #1506586.

## References

- [1] S. K. Bajracharya and C. V. Lopes. Mining search topics from a code search engine usage log. In *MSR*, pages 111–120. Citeseer, 2009.
- [2] A. Barua, S. W. Thomas, and A. E. Hassan. What are developers talking about? an analysis of topics and trends in stack overflow. *Empirical Software Engineering*, 19(3):619–654, 2014.
- [3] R. P. Beausoleil. MOSS multiobjective scatter search applied to non-linear multiple criteria optimization. *European Journal of Operational Research*, 169(2):426–449, 2006.
- [4] J. Bergstra and Y. Bengio. Random search for hyper-parameter optimization. *Journal of Machine Learning Research*, 13(Feb):281–305, 2012.
- [5] D. Binkley, D. Heinz, D. Lawrie, and J. Overfelt. Understanding lda in source code analysis. In *Proceedings of the 22nd International Conference on Program Comprehension*, pages 26–36. ACM, 2014.
- [6] S. Bird. Nltk: the natural language toolkit. In *Proceedings of the COLING/ACL on Interactive presentation sessions*, pages 69–72. Association for Computational Linguistics, 2006.
- [7] D. M. Blei, A. Y. Ng, and M. I. Jordan. Latent dirichlet allocation. In *The Journal of machine Learning research*, volume 3, pages 993–1022. JMLR.org, 2003.
- [8] L. Breiman. Random forests. *Machine Learning*, 45(1):5–32, 2001.
- [9] T.-H. Chen, W. Shang, M. Nagappan, A. E. Hassan, and S. W. Thomas. Topic-based software defect explanation. *Journal of Systems and Software*, 2016.
- [10] T.-H. Chen, S. W. Thomas, M. Nagappan, and A. E. Hassan. Explaining software defects using topic models. In *Mining Software Repositories (MSR), 2012 9th IEEE Working Conference on*, pages 189–198. IEEE, 2012.
- [11] I. Chiha, J. Ghabi, and N. Liouane. Tuning pid controller with multi-objective differential evolution. In *Communications Control and Signal Processing (ISCCSP), 2012 5th International Symposium on*, pages 1–4. IEEE, 2012.
- [12] K. O. Elish and M. O. Elish. Predicting defect-prone software modules using support vector machines. *Journal of Systems and Software*, 81(5):649–660, 2008.
- [13] M. S. Feather and T. Menzies. Converging on the optimal attainment of requirements. In *Requirements Engineering, 2002. Proceedings. IEEE Joint International Conference on*, pages 263–270. IEEE, 2002.
- [14] R.-S. Feldman. *J, The Text Mining Handbook*. New York: Cambridge University Press, 2006.
- [15] W. Fu, T. Menzies, and X. Shen. Tuning for software analytics: Is it really necessary? *Information and Software Technology*, 76:135–146, 2016.
- [16] Y. Fu, M. Yan, X. Zhang, L. Xu, D. Yang, and J. D. Kymer. Automated classification of software change messages by semi-supervised latent dirichlet allocation. *Information and Software Technology*, 57:369–377, 2015.
- [17] L. V. Galvis Carreño and K. Winbladh. Analysis of user comments: an approach for software requirements evolution. In *Proceedings of the 2013 International Conference on Software Engineering*, pages 582–591. IEEE Press, 2013.
- [18] V. Garousi and M. V. Mäntylä. Citations, research topics and active countries in software engineering: A bibliometrics study. *Computer Science Review*, 19:56–77, 2016.
- [19] J. H. Gennari, P. Langley, and D. Fisher. Models of incremental concept formation. *Artificial Intelligence*, 40(1):11 – 61, 1989.
- [20] F. Glover and C. McMillan. The general employee scheduling problem. an integration of ms and ai. *Computers & operations research*, 13(5):563–573, 1986.
- [21] A. T. Goldberg. *On the complexity of the satisfiability problem*. PhD thesis, New York University, 1979.
- [22] S. Grant, J. R. Cordy, and D. B. Skillicorn. Using heuristics to estimate an appropriate number of latent topics in source code analysis. *Science of Computer Programming*, 78(9):1663–1678, 2013.
- [23] T. L. Griffiths and M. Steyvers. Finding scientific topics. *Proceedings of the National academy of Sciences*, 101(suppl 1):5228–5235, 2004.
- [24] E. Guzman and W. Maalej. How do users like this feature? a fine grained sentiment analysis of app reviews. In *Requirements Engineering Conference (RE), 2014 IEEE 22nd International*, pages 153–162. IEEE, 2014.
- [25] A. Hindle, C. Bird, T. Zimmermann, and N. Nagappan. Relating requirements to implementation via topic analysis: Do topics extracted from requirements make sense to managers and developers? In *Software Maintenance (ICSM), 2012 28th IEEE International Conference on*, pages 243–252. IEEE, 2012.
- [26] A. Hindle, N. A. Ernst, M. W. Godfrey, and J. Mylopoulos. Automated topic naming to support cross-project analysis of software maintenance activities. In *Proceedings of the 8th Working Conference on Mining Software Repositories*, pages 163–172. ACM, 2011.
- [27] S. Koltcov, O. Koltsova, and S. Nikolenko. Latent dirichlet allocation: stability and applications to studies of user-generated content. In *Proceedings of the 2014 ACM conference on Web science*, pages 161–165. ACM, 2014.
- [28] J. Krall, T. Menzies, and M. Davies. Gale: Geometric active learning for search-based software engineering. *Software Engineering, IEEE Transactions on*, 41(10):1001–1018, 2015.
- [29] R. Krishna, Z. Yu, A. Agrawal, M. Dominguez, and D. Wolf. The bigse project: lessons learned from validating industrial text mining. In *Proceedings of the 2nd International Workshop on BIG Data Software Engineering*, pages 65–71. ACM, 2016.
- [30] J. H. Lau, D. Newman, and T. Baldwin. Machine reading tea leaves: Automatically evaluating topic coherence and topic model quality. In *EACL*, pages 530–539, 2014.
- [31] L. Layman, 2016.
- [32] L. Layman, A. P. Nikora, J. Meek, and T. Menzies. Topic modeling of nasa space system problem reports: research in practice. In *Proceedings of the 13th International Workshop on Mining Software Repositories*, pages 303–314. ACM, 2016.
- [33] T.-D. B. Le, F. Thung, and D. Lo. Predicting effectiveness of ir-based bug localization techniques. In *2014 IEEE 25th International Symposium on Software Reliability Engineering*, pages 335–345. IEEE, 2014.
- [34] S. Lessmann, B. Baesens, C. Mues, and S. Pietsch. Benchmarking classification models for software defect prediction: A proposed framework and novel findings. *IEEE Transactions on Software Engineering*, 34(4):485–496, July 2008.
- [35] M. Linares-Vásquez, B. Dit, and D. Poshvanyk. An exploratory analysis of mobile development issues using stack overflow. In *Proceedings of the 10th Working Conference on Mining Software Repositories*, pages 93–96. IEEE Press, 2013.
- [36] S. Lohar, S. Amornborvornwong, A. Zisman, and J. Cleland-Huang. Improving trace accuracy through data-driven configuration and composition of tracing features. In *Proceedings of the 2013 9th Joint Meeting on Foundations of Software Engineering*, pages 378–388. ACM, 2013.



- [37] S. K. Lukins, N. A. Kraft, and L. H. Etzkorn. Bug localization using latent dirichlet allocation. *Information and Software Technology*, 52(9):972–990, 2010.
- [38] W. Martin, M. Harman, Y. Jia, F. Sarro, and Y. Zhang. The app sampling problem for app store mining. In *Mining Software Repositories (MSR), 2015 IEEE/ACM 12th Working Conference on*, pages 123–133. IEEE, 2015.
- [39] G. Maskeri, S. Sarkar, and K. Heafield. Mining business topics in source code using latent dirichlet allocation. In *Proceedings of the 1st India software engineering conference*, pages 113–120. ACM, 2008.
- [40] X. Meng, J. Bradley, B. Yuvaz, E. Sparks, S. Venkataraman, D. Liu, J. Freeman, D. Tsai, M. Amde, S. Owen, et al. Millib: Machine learning in apache spark. *JMLR*, 17(34):1–7, 2016.
- [41] T. Menzies. Improving iv&v techniques through the analysis of project anomalies: Text mining pits issue reports-final report. *Citeseer*, 2008.
- [42] T. Menzies, O. Elrawas, J. Hihn, M. Feather, R. Madachy, and B. Boehm. The business case for automated software engineering. In *Proceedings of the twenty-second IEEE/ACM international conference on Automated software engineering*, pages 303–312. ACM, 2007.
- [43] T. Menzies, R. Krishna, and D. Pryor. The Promise Repository of Empirical Software Engineering Data. <http://openscience.us/repo>, 2015.
- [44] T. Menzies and A. Marcus. Automated severity assessment of software defect reports. In *Software Maintenance, 2008. ICSM 2008. IEEE International Conference on*, pages 346–355. IEEE, 2008.
- [45] G. A. Miller. The magical number seven, plus or minus two: Some limits on our capacity for processing information. *Psychological review*, 63(2):81, 1956.
- [46] T. Minka and J. Lafferty. Expectation-propagation for the generative aspect model. In *Proceedings of the Eighteenth conference on Uncertainty in artificial intelligence*, pages 352–359. Morgan Kaufmann Publishers Inc., 2002.
- [47] J. Molina, M. Laguna, R. Martí, and R. Caballero. Sspmo: A scatter tabu search procedure for non-linear multiobjective optimization. *INFORMS Journal on Computing*, 19(1):91–100, 2007.
- [48] A. Nadkarni and N. Yezhkova. Structured versus unstructured data: The balance of power continues to shift. *IDC (Industry Development and Models) Mar*, 2014.
- [49] A. J. Nebro, F. Luna, E. Alba, B. Dorronsoro, J. J. Durillo, and A. Beham. Abyss: Adapting scatter search to multiobjective optimization. *Evolutionary Computation, IEEE Transactions on*, 12(4):439–457, 2008.
- [50] S. I. Nikolenko, S. Koltcov, and O. Koltsova. Topic modelling for qualitative studies. *Journal of Information Science*, page 0165551515617393, 2015.
- [51] R. Oliveto, M. Gethers, D. Poshyvanyk, and A. De Lucia. On the equivalence of information retrieval methods for automated traceability link recovery. In *Program Comprehension (ICPC), 2010 IEEE 18th International Conference on*, pages 68–71. IEEE, 2010.
- [52] M. G. Omran, A. P. Engelbrecht, and A. Salman. Differential evolution methods for unsupervised image classification. In *Evolutionary Computation, 2005. The 2005 IEEE Congress on*, volume 2, pages 966–973. IEEE, 2005.
- [53] D. OCallaghan, D. Greene, J. Carthy, and P. Cunningham. An analysis of the coherence of descriptors in topic modeling. *Expert Systems with Applications*, 42(13):5645–5657, 2015.
- [54] H. Pan, M. Zheng, and X. Han. Particle swarm-simulated annealing fusion algorithm and its application in function optimization. In *Computer Science and Software Engineering, 2008 International Conference on*, volume 1, pages 78–81. IEEE, 2008.
- [55] A. Panichella, B. Dit, R. Oliveto, M. Di Penta, D. Poshyvanyk, and A. De Lucia. How to effectively use topic models for software engineering tasks? an approach based on genetic algorithms. In *Proceedings of the 2013 International Conference on Software Engineering*, pages 522–531. IEEE Press, 2013.
- [56] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, et al. Scikit-learn: Machine learning in python. *Journal of Machine Learning Research*, 12(Oct):2825–2830, 2011.
- [57] M. Porter. The Porter Stemming Algorithm, 1980.
- [58] D. M. Powers. Evaluation: from precision, recall and f-measure to roc, informedness, markedness and correlation. 2011.
- [59] S. Rao and A. Kak. Retrieval from software libraries for bug localization: a comparative study of generic and composite text models. In *Proceedings of the 8th Working Conference on Mining Software Repositories*, pages 43–52. ACM, 2011.
- [60] A. S. Sayyad, J. Ingram, T. Menzies, and H. Ammar. Scalable product line configuration: A straw to break the camel’s back. In *Automated Software Engineering (ASE), 2013 IEEE/ACM 28th International Conference on*, pages 465–474. IEEE, 2013.
- [61] R. Storn and K. Price. Differential evolution—a simple and efficient heuristic for global optimization over continuous spaces. *Journal of global optimization*, 11(4):341–359, 1997.
- [62] X. Sun, B. Li, H. Leung, B. Li, and Y. Li. Msr4sm: Using topic models to effectively mining software repositories for software maintenance tasks. *Information and Software Technology*, 66:1–12, 2015.
- [63] C. Tantithamthavorn, S. McIntosh, A. E. Hassan, and K. Matsumoto. Automated Parameter Optimization of Classification techniques for Defect Prediction Models. In *Proc. of the International Conference on Software Engineering (ICSE)*, pages 321–332, 2016.
- [64] S. W. Thomas. Mining software repositories using topic models. In *Proceedings of the 33rd International Conference on Software Engineering*, pages 1138–1139. ACM, 2011.
- [65] S. W. Thomas, B. Adams, A. E. Hassan, and D. Blostein. Studying software evolution using topic models. *Science of Computer Programming*, 80:457–479, 2014.
- [66] S. W. Thomas, H. Hemmati, A. E. Hassan, and D. Blostein. Static test case prioritization using topic models. *Empirical Software Engineering*, 19(1):182–212, 2014.
- [67] K. Tian, M. Revelle, and D. Poshyvanyk. Using latent dirichlet allocation for automatic categorization of software. In *Mining Software Repositories, 2009. MSR’09. 6th IEEE International Working Conference on*, pages 163–166. IEEE, 2009.
- [68] B. Vasilescu, A. Serebrenik, and T. Mens. A historical dataset of software engineering conferences. In *Proceedings of the 10th Working Conference on Mining Software Repositories*, pages 373–376. IEEE Press, 2013.
- [69] J. Vesterström and R. Thomsen. A comparative study of differential evolution, particle swarm optimization, and evolutionary algorithms on numerical benchmark problems. In *Evolutionary Computation, 2004. CEC2004. Congress on*, volume 2, pages 1980–1987. IEEE, 2004.
- [70] T. Wang, M. Harman, Y. Jia, and J. Krinke. Searching for better configurations: a rigorous approach to clone evaluation. In *Proceedings of the 2013 9th Joint Meeting on Foundations of Software Engineering*, pages 455–465. ACM, 2013.
- [71] X. Wei and W. B. Croft. Lda-based document models for ad-hoc retrieval. In *Proceedings of the 29th annual international ACM SIGIR conference on Research and development in information retrieval*, pages 178–185. ACM, 2006.
- [72] Y. Yang. *Improving the Usability of Topic Models*. PhD thesis, NORTH-WESTERN UNIVERSITY, 2015.
- [73] Y. Yang, Y. Zhou, J. Liu, Y. Zhao, H. Lu, L. Xu, B. Xu, and H. Leung. Effort-aware just-in-time defect prediction: Simple unsupervised models could be better than supervised models. In *Proceedings of the 2016 24th ACM SIGSOFT International Symposium on Foundations of Software Engineering, FSE 2016*, pages 157–168, New York, NY, USA, 2016. ACM.
- [74] W. Zhao, J. J. Chen, R. Perkins, Z. Liu, W. Ge, Y. Ding, and W. Zou. A heuristic approach to determine an appropriate number of topics in topic modeling. *BMC bioinformatics*, 16(13):1, 2015.
- [75] M. Zuluaga, G. Sergeant, A. Krause, and M. Püschel. Active learning for multi-objective optimization. In *Proceedings of the 30th International Conference on Machine Learning*, pages 462–470, 2013.